


Connecting to SQL Databases with Python

11 April 2023



Day One Contents

kubrick

Connecting to Databases from Python

Configuring PyODBC

Querying Data & Processing Results

Working with Data Types

Connecting to Databases from Python

Database API Specification

Python has a well-established standard pattern for accessing databases - the DBI specification, currently in its v2.0 release
See <https://peps.python.org/pep-0249/>

DBI came out of PEP-249 and specifies what you should expect in an API for accessing databases

The standard is very useful as it allows programmers to understand the features of a database driver quickly and become productive with it, so long as it adheres to the DBI Specifications

ODBC Open Database Connectivity

ODBC was established as a standard for Database Connectivity in the 1990's
(way before Python was developed)

The Standard was proposed from work done at Microsoft and Samba and widely adopted by other database vendors

The Standard defines how to connect to a database and how data will be passed between applications (clients) and databases

ODBC was initially available on the Windows platform, but it was soon ported to Mac OSX and Linux

Almost every database vendor or developer will have created an ODBC driver for their database
The ODBC driver is supplied as a piece of executable code that is installed on the operating system
ODBC drivers operate at quite a low level and are generally written in C or C++
They provide the API that higher level software interacts with.

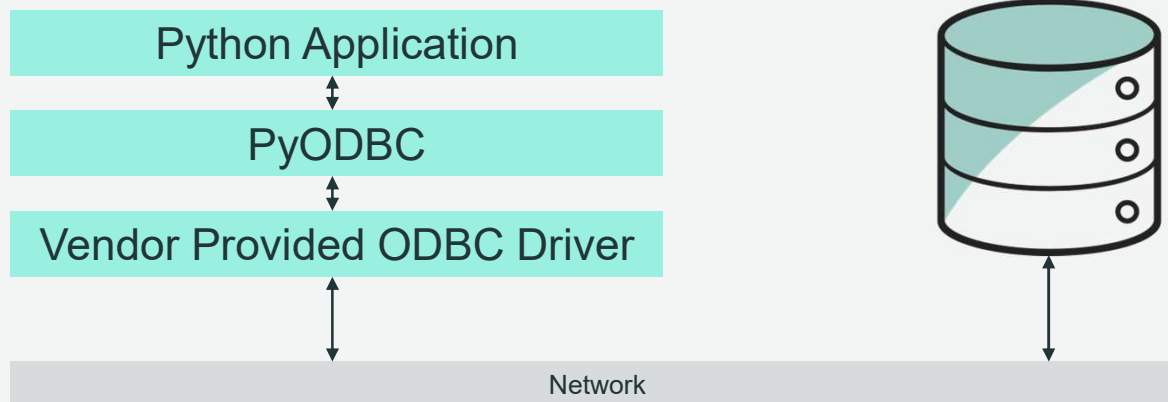
ODBC drivers are available for Microsoft SQL Server, Oracle, MariaDB, PostgreSQL, SQLite and many other databases and data sources (such as Excel)

PyODBC – Python's DBI V2 Interface to ODBC

PyODBC provides an implementation of the Python DBI V2.0 interface on top of a vendor provided ODBC Driver

It provides Python programmers with the familiar DBI API and provides a useful degree of insulation from the differences in the underlying database technologies

pyodbc has a splendid support wiki at github.com/mkleehammer/pyodbc/wiki and is essential post course reading



PyODBC – Benefits for Python Programmers

Operating System
Independent
Same Python code runs on
Windows, Mac, Linux

Database Independent
But there are still some
differences in the way
databases operate

Simple and familiar
programming interface from
DBI V2.0

Setting Up the PyODBC Development Environment

The set up is designed to work on a Windows 10 PC

We will be connecting to 2 different databases in these sessions

- a SQL Server Database you accessed earlier in your SQL training
- a local SQLite file-based database
- It will help to install SQL Server Management Studio for looking at SQL objects
<https://docs.microsoft.com/en-us/sql/ssms/download-sql-server-management-studio-ssms?view=sql-server-ver16>
- The ODBC driver for Microsoft SQL Server is available at :
<https://docs.microsoft.com/en-us/sql/connect/odbc/download-odbc-driver-for-sql-server?view=sql-server-ver16>
- Anaconda Python – should already have pyodbc and sqlite installed
use “conda install” to install if not
- If using another Python distribution create a virtual env and install the PyODBC module with pip
- Visual Studio Code with the Python extensions is my preferred development IDE, but you can use a the Jupyter notebook server that comes with Anaconda

```
Anaconda Prompt (Anaconda3)
(base) C:\Users\DavidBurnham>conda list odbc
# packages in environment at C:\tools\Anaconda3:
#
# Name          Version          Build      Channel
pyodbc          4.0.31           py39hd77b12b_0

(base) C:\Users\DavidBurnham>conda list sqlite
# packages in environment at C:\tools\Anaconda3:
#
# Name          Version          Build      Channel
sqlite          3.36.0           h2bbff1b_0

(base) C:\Users\DavidBurnham>
```

PyODBC Development Environment

Essential:

Python 3.7+ (Anaconda works fine or you can install vanilla Python and install the additional modules with pip3)

PyODBC & Sqlite3 python modules

Microsoft SQL Server ODBC Driver (V17+) – download from Microsoft

(<https://docs.microsoft.com/en-us/sql/connect/odbc/download-odbc-driver-for-sql-server?view=sql-server-ver16>)

Development IDE – one that you are productive with Visual Studio Code is my preference

Nice to have:

Microsoft SQL Management Studio– download from Microsoft

(<https://docs.microsoft.com/en-us/sql/ssms/download-sql-server-management-studio-ssms?view=sql-server-ver16>)

PyODBC Object Model – Module Level




Configuration
attributes



`pyodbc.connect()`

Once the `pyodbc.connect()` function has been called successfully a connection object is created. The connection object has a method (`cursor`) that returns a cursor allowing the application to submit SQL requests to the database.

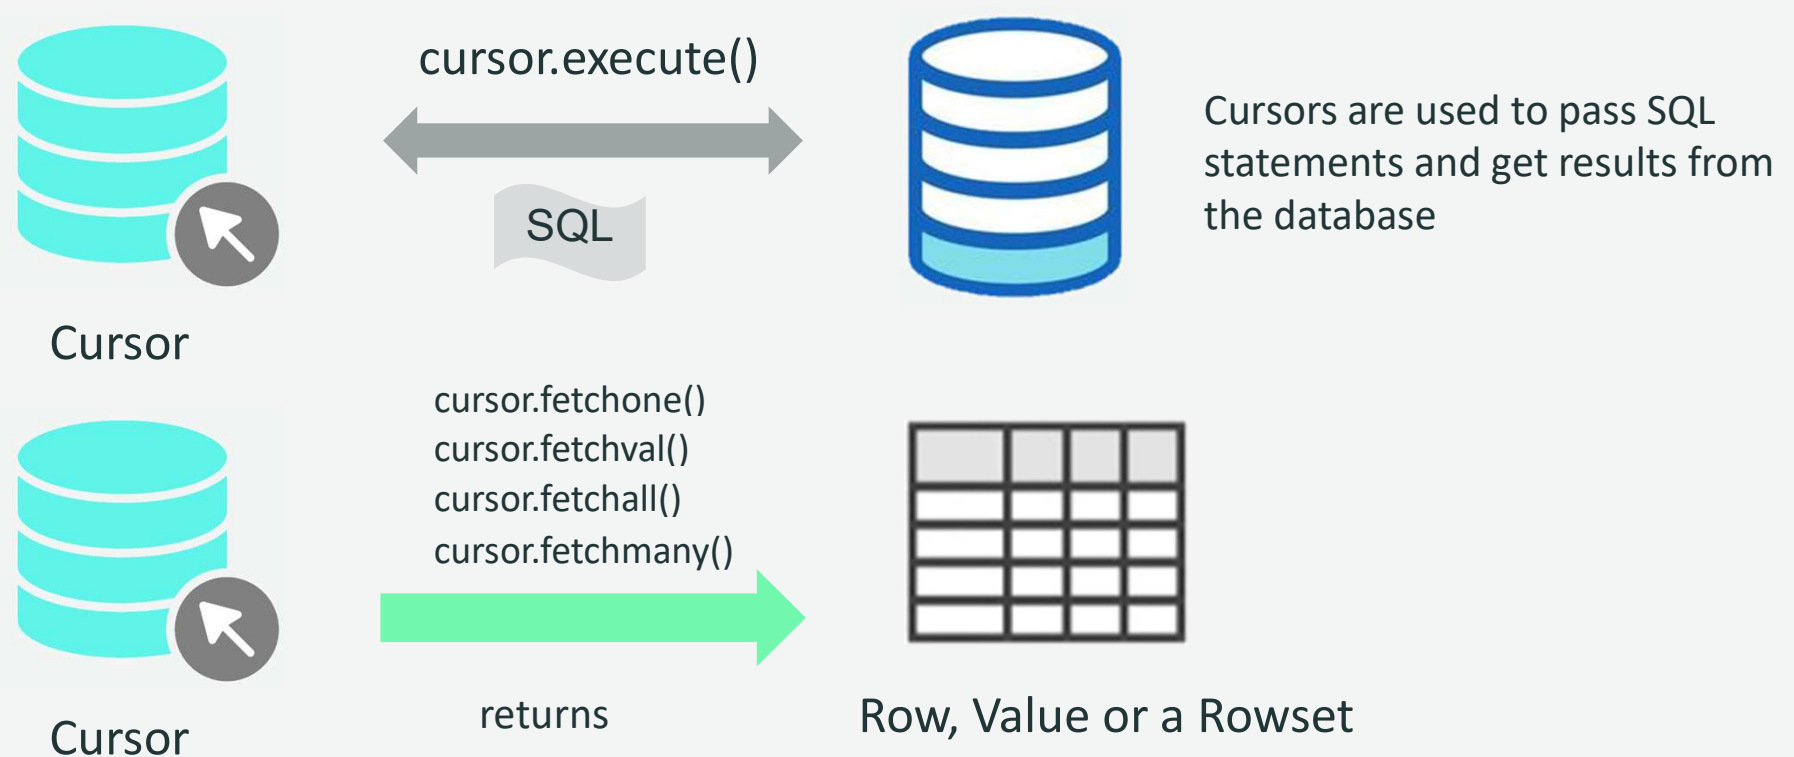
PyODBC Object Model – connect returns connection kubrick



Once the `pyodbc.connect()` function has been called successfully a connection object is created.

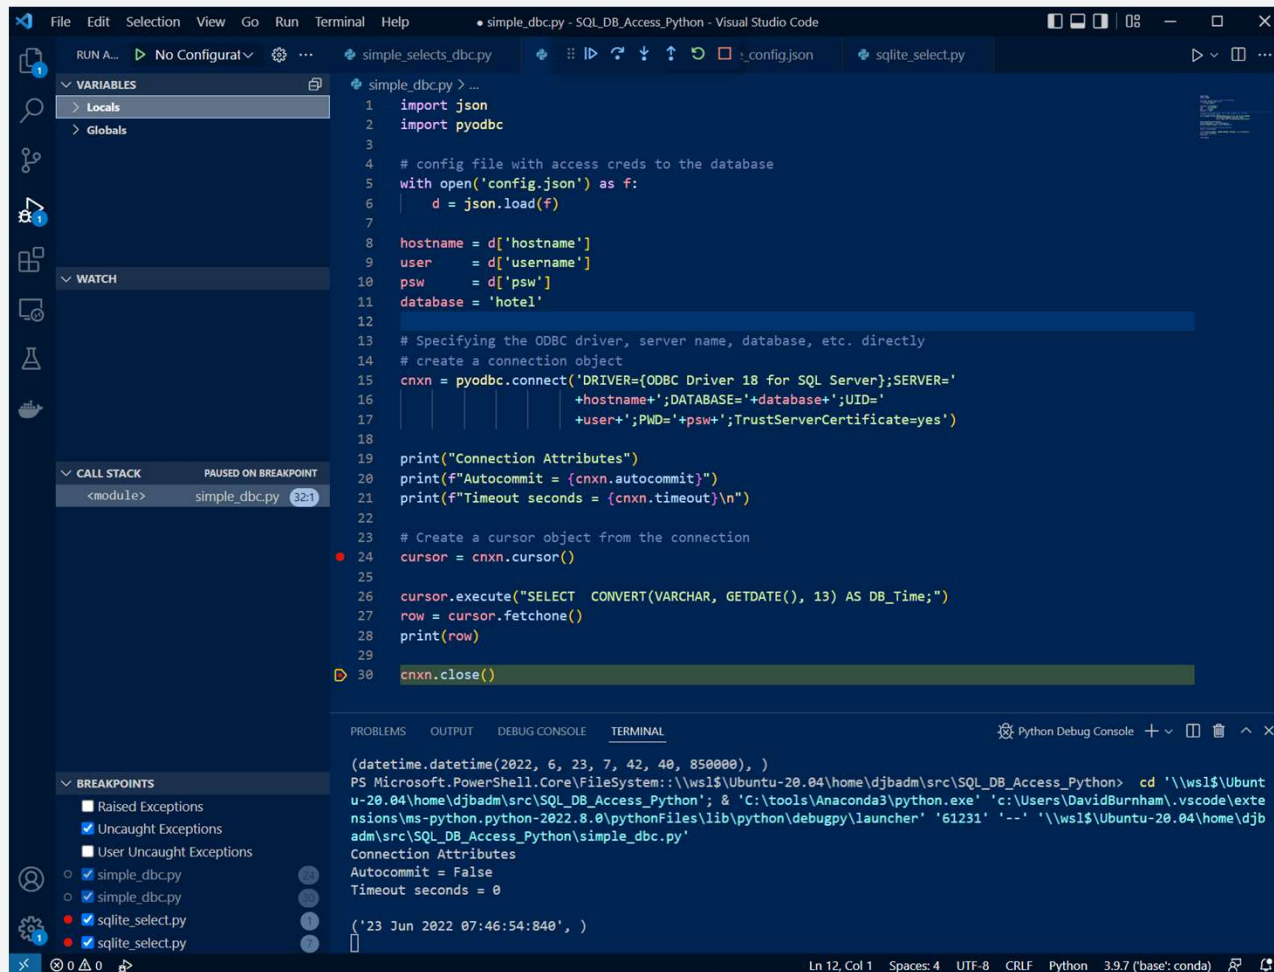
The connection object has a method (`cursor`) that returns a cursor allowing the application to submit SQL requests to the database.

PyODBC Object Model – Cursor Object



Cursors are python representations of a database cursor – they may or may not be implemented on top of database cursors. They are not identical.

Demonstration – Connecting to a Database



```
File Edit Selection View Go Run Terminal Help • simple_dbc.py - SQL_DB_Access_Python - Visual Studio Code
RUN A... No Configurat... simple_selects_dbc.py _config.json sqlite_select.py
VARIABLES
> Locals
> Globals
WATCH
CALL STACK PAUSED ON BREAKPOINT
<module> simple_dbc.py 32:1
BREAKPOINTS
[ ] Raised Exceptions
[x] Uncaught Exceptions
[ ] User Uncaught Exceptions
[x] simple_dbc.py
[x] simple_dbc.py
[x] sqlite_select.py
[x] sqlite_select.py
1 import json
2 import pyodbc
3
4 # config file with access creds to the database
5 with open('config.json') as f:
6     d = json.load(f)
7
8 hostname = d['hostname']
9 user = d['username']
10 psw = d['psw']
11 database = 'hotel'
12
13 # Specifying the ODBC driver, server name, database, etc. directly
14 # create a connection object
15 cnxn = pyodbc.connect('DRIVER={ODBC Driver 18 for SQL Server};SERVER='
16                       +hostname+';DATABASE='+database+';UID='
17                       +user+';PWD='+psw+';TrustServerCertificate=yes')
18
19 print("Connection Attributes")
20 print(f"Autocommit = {cnxn.autocommit}")
21 print(f"Timeout seconds = {cnxn.timeout}\n")
22
23 # Create a cursor object from the connection
24 cursor = cnxn.cursor()
25
26 cursor.execute("SELECT CONVERT(VARCHAR, GETDATE(), 13) AS DB_Time;")
27 row = cursor.fetchone()
28 print(row)
29
30 cnxn.close()
```

Python Debug Console

```
(datetime.datetime(2022, 6, 23, 7, 42, 40, 850000), )
PS Microsoft.PowerShell.Core\FileSystem::\\ws1$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python> cd '\\ws1$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python'; & 'C:\\tools\\Anaconda3\\python.exe' 'c:\\Users\\DavidBurnham\\.vscode\\extensions\\ms-python.python-2022.8.0\\pythonFiles\\lib\\python\\debugpy\\launcher' '61231' '--' '\\ws1$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python\\simple_dbc.py'
Connection Attributes
Autocommit = False
Timeout seconds = 0

('23 Jun 2022 07:46:54:840', )
```

Exercise Ex1

Make sure your PC Development environment is set up to develop with Python PyODBC

Connect to the database we used for the SQL training

Here is a json fragment for your config.json file :

```
{  "hostname": "ec2-35-176-213-89.eu-west-2.compute.amazonaws.com",
  "database": "hotel",
  "username": "YourUserName",
  "psw": "*****"
}
```

Identify the Operating System the database is running on :

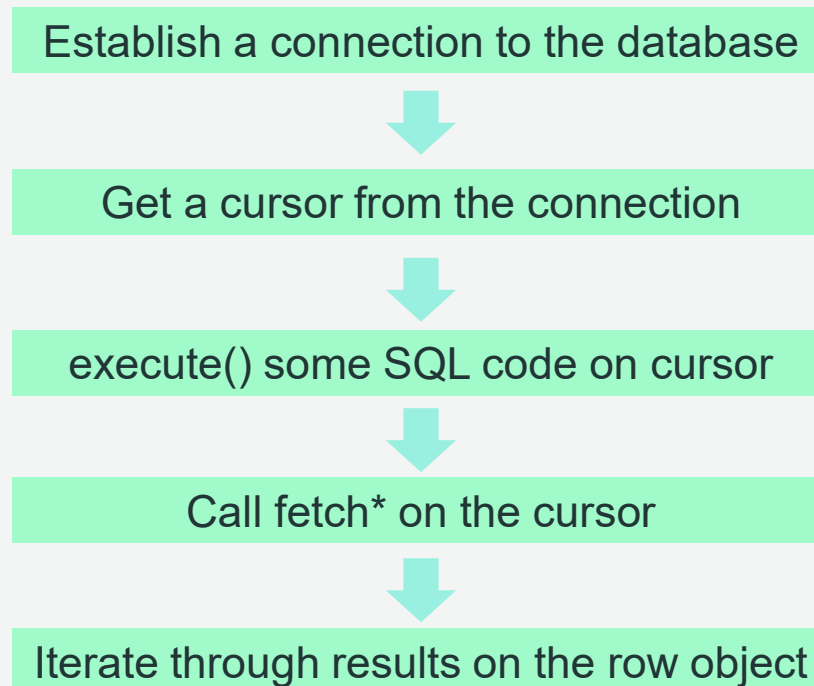
Hint look up the information stored in the global variable @@VERSION

Work as individuals (you need to all set up a working pydbc development environment)

The exercise should take about 20 minutes maximum

As always message me if you are having difficulties ... 😊

Basic Query Flow



Points of note about pyodbc cursor metadata

After a piece of SQL has been executed on a pyodbc cursor – the cursor contains useful metadata
The **cursor.description** attribute is populated with the column name and datatype for each row being returned - the attribute is a tuple of tuples

cursor.messages contains any messages from the database that may have resulted from the cursor execution

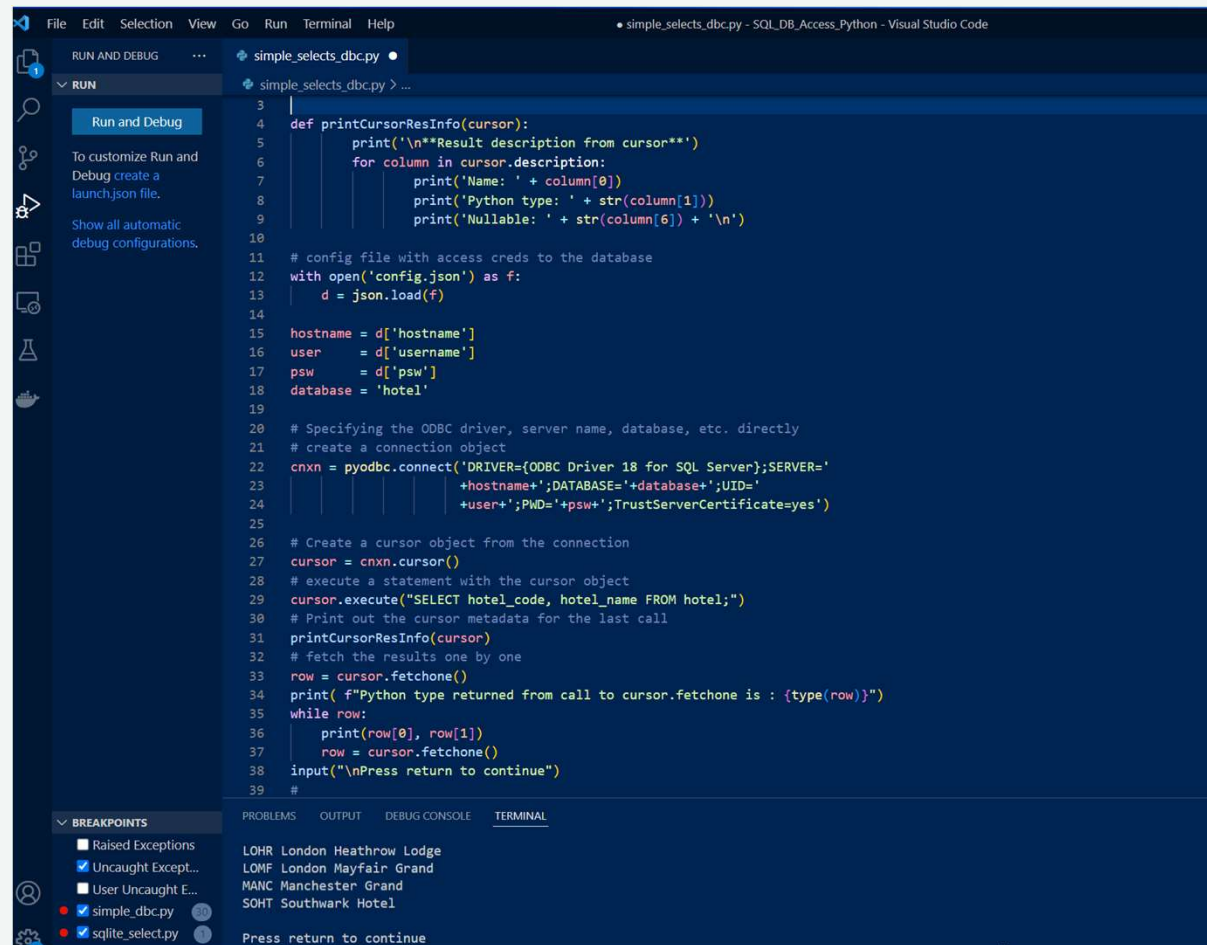
cursor.rowcount contains the number of rows that have been affected by any DML statements.



Its implementation is database driver specific and select statements (which don't change any rows) just return a cursor.rowcount of -1 for the SQL Server ODBC driver

See <https://github.com/mkleehammer/pyodbc/wiki/Cursor> for further details

Demo – Executing Select Queries



```
File Edit Selection View Go Run Terminal Help
simple_selects_dbc.py - SQL_DB_Access_Python - Visual Studio Code

RUN AND DEBUG ...
simple_selects_dbc.py > ...

Run and Debug
To customize Run and Debug create a launch.json file.
Show all automatic debug configurations.

3
4 def printCursorResInfo(cursor):
5     print('\n**Result description from cursor**')
6     for column in cursor.description:
7         print('Name: ' + column[0])
8         print('Python type: ' + str(column[1]))
9         print('Nullable: ' + str(column[6]) + '\n')
10
11 # config file with access creds to the database
12 with open('config.json') as f:
13     d = json.load(f)
14
15     hostname = d['hostname']
16     user = d['username']
17     psd = d['psw']
18     database = 'hotel'
19
20 # Specifying the ODBC driver, server name, database, etc. directly
21 # create a connection object
22 cnxn = pyodbc.connect('DRIVER={ODBC Driver 18 for SQL Server};SERVER='
23                      +hostname+';DATABASE='+database+';UID='
24                      +user+';PWD='+psw+';TrustServerCertificate=yes')
25
26 # Create a cursor object from the connection
27 cursor = cnxn.cursor()
28 # execute a statement with the cursor object
29 cursor.execute("SELECT hotel_code, hotel_name FROM hotel;")
30 # Print out the cursor metadata for the last call
31 printCursorResInfo(cursor)
32 # fetch the results one by one
33 row = cursor.fetchone()
34 print( f"Python type returned from call to cursor.fetchone is : {type(row)}")
35 while row:
36     print(row[0], row[1])
37     row = cursor.fetchone()
38 input("\nPress return to continue")
39 #

BREAKPOINTS
[ ] Raised Exceptions
[ ] Uncaught Except...
[ ] User Uncaught E...
[ ] simple_dbc.py
[ ] sqlite_select.py

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
LOHR London Heathrow Lodge
LOMF London Mayfair Grand
MANC Manchester Grand
SOHT Southwark Hotel
Press return to continue
```


Using parameters with execute() in SELECTs

Parameters can be passed to a SQL statement by including a ? character in the text of the SQL and passing the values to substitute for the ? character as parameters.

Parameters can also be used with any other SQL statements such as inserts, updates, deleted and stored procedure executions. This is extremely useful as it ensures you do not have to build SQL statements using string concatenation which is a major cause of SQL Injection vulnerabilities in applications.

There is an issue with the pyodbc module and the SQL Server driver you need to be aware of though. The pyodbc driver tries to optimize SQL varchar variable sizes and will parse each python string variable and create a new (different) SQL submission to the driver for each different string length.

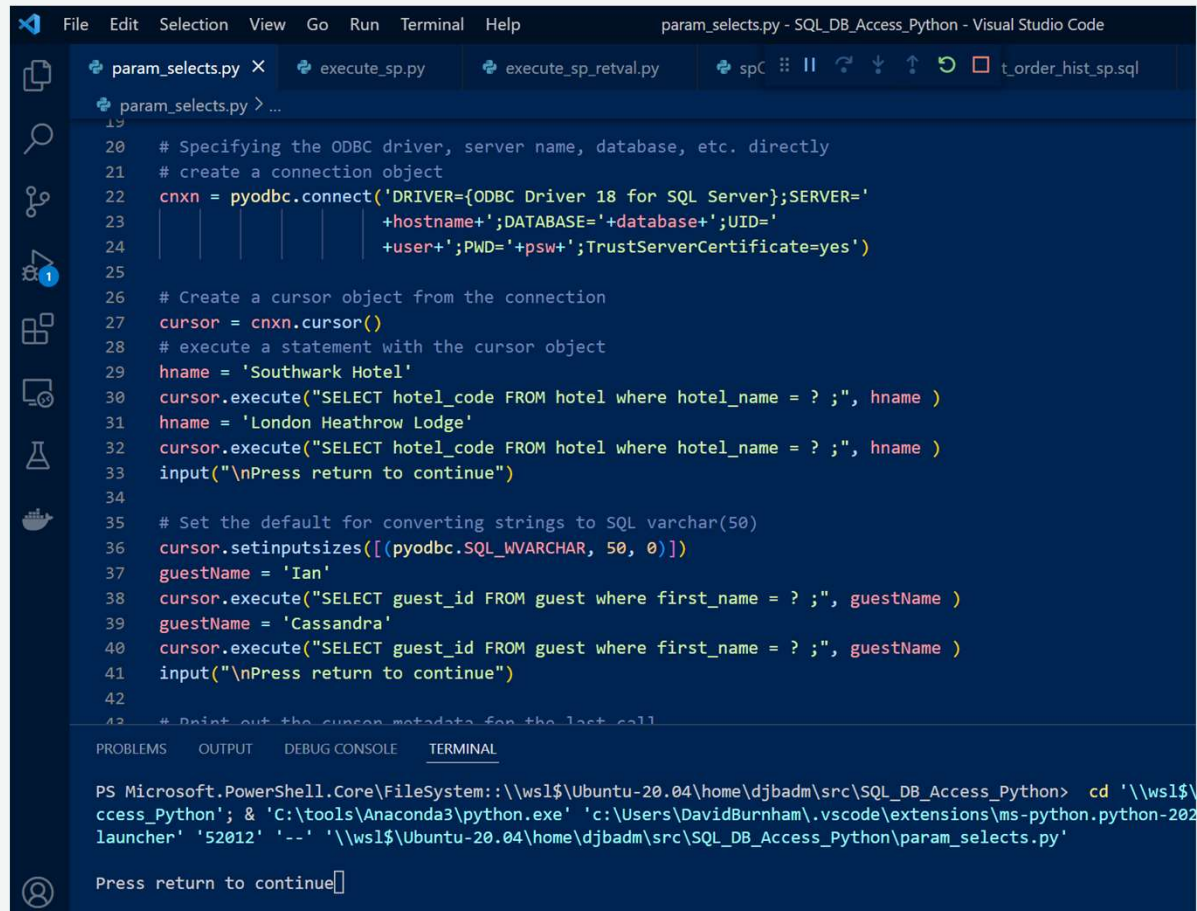


This is very bad for database optimisation as it results in many SQL execution plans being generated for what is essentially the same SQL statement.

We can avoid this by explicitly setting the size of strings that pyodbc will use for strings :

```
cursor.setinputsizes(pyodbc.SQL_WVARCHAR, 50 ,0) # sets strings to be converted to varchar(50)
```

Demo – Executing Select Queries with Parameters



The screenshot shows the Visual Studio Code editor with a file named `param_selects.py` open. The script is a Python program that uses the `pyodbc` library to connect to a SQL Server database and execute queries with parameters. The script includes comments and uses `input()` for user interaction. The terminal at the bottom shows the command used to run the script in a Windows PowerShell environment.

```
19
20 # Specifying the ODBC driver, server name, database, etc. directly
21 # create a connection object
22 cnxn = pyodbc.connect('DRIVER={ODBC Driver 18 for SQL Server};SERVER='
23                       +hostname+';DATABASE='+database+';UID='
24                       +user+';PWD='+psw+';TrustServerCertificate=yes')
25
26 # Create a cursor object from the connection
27 cursor = cnxn.cursor()
28 # execute a statement with the cursor object
29 hname = 'Southwark Hotel'
30 cursor.execute("SELECT hotel_code FROM hotel where hotel_name = ? ;", hname )
31 hname = 'London Heathrow Lodge'
32 cursor.execute("SELECT hotel_code FROM hotel where hotel_name = ? ;", hname )
33 input("\nPress return to continue")
34
35 # Set the default for converting strings to SQL varchar(50)
36 cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
37 guestName = 'Ian'
38 cursor.execute("SELECT guest_id FROM guest where first_name = ? ;", guestName )
39 guestName = 'Cassandra'
40 cursor.execute("SELECT guest_id FROM guest where first_name = ? ;", guestName )
41 input("\nPress return to continue")
42
43 # Print out the cursor metadata for the last call
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
PS Microsoft.PowerShell.Core\FileSystem::\\wsl$Ubuntu-20.04\home\djbadm\src\SQL_DB_Access_Python> cd '\\wsl$\\
ccess_Python'; & 'C:\tools\Anaconda3\python.exe' 'c:\Users\DavidBurnham\.vscode\extensions\ms-python.python-202
launcher' '52012' '--' '\\wsl$Ubuntu-20.04\home\djbadm\src\SQL_DB_Access_Python\param_selects.py'

Press return to continue
```

Inserts updates deletes using pyodbc

SQL INSERT, UPDATE and DELETE statements can all be called and parameterised in the same way that SELECT statements can be. In the next demonstration we will see how this can be performed.

For the purposes of the demonstration we will set the connection.autocommit parameter to True, which will cause the change to be applied to the database after each statement is executed. This is not the default, but we will discuss how this parameter is managed later in the course.

Demo – Inserts Updates & Deletes with Parameters

```
File Edit Selection View Go Run Terminal Help insert_upd_del.py - SQL_DB_Access_Python - Visual Studio Code
insert_upd_del.py X
insert_upd_del.py > ...
30 # Create a cursor object from the connection
31 cursor = cnxn.cursor()
32 cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
33
34 Sql = '''SELECT car_id, manufacturer, model, fuel, engine_capacity,
35 engine_power FROM [dbo].[cars]; '''
36
37 cursor.execute(Sql)
38
39 # Print out the cursor metadata for the last call
40 printCursorResInfo(cursor)
41
42 allRows = cursor.fetchall()
43 for row in allRows:
44     print(row)
45
46 input("\nPress return to continue")
47
48 Sql = '''INSERT INTO [dbo].[cars]
49 (car_id, manufacturer, model, fuel, engine_capacity, engine_power)
50 VALUES (?, ?, ?, ?, ?, ?) ; '''
51
52 newCar1 = (8, 'Skoda', 'Submariner', 'Petrol', 2498, 220)
53 newCar2 = (9, 'Saab', 'Talladega', 'Petrol', 1948, 197)
54
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
PS Microsoft.PowerShell.Core\FileSystem::\\wsl$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python>
```

```
SQLQuery15.sql - ec2-35-176-213-89.eu-west-2.compute.amazonaws.com:djbpyodbc (trainer (73))* - Microsoft SQL Server
File Edit View Query Project Tools Window Help
djbpodbc Execute
SQLQuery15.sql - ec...odbc (trainer (73))* SQLQuery14.sql - ec...odbc (trainer (67))* SQLQuery13.sql - ec...odbc (
/***** Script for SelectTopNRows command from SSMS *****/
SELECT TOP (1000) [car_id]
, [manufacturer]
, [model]
, [fuel]
, [engine_capacity]
, [engine_power]
FROM [djbpodbc].[dbo].[cars]

100 %
Results Messages
car_id manufacturer model fuel engine_capacity engine_power
1 1 Ford Topix Petrol 1298 80
2 2 Vauxhall Sparkler Electricity NULL 120
3 3 Nissan Banjo Diesel 2198 177
4 4 Tesla Zappy Electrickery NULL 212
5 5 Land Rover Lemon Diesel 2488 149
6 6 Seat Massive Electricity NULL 177
7 7 VW Rolf Electricity NULL 120
```

Exercise Ex2

Selecting inserting and updating with parameters

Do this exercise in your pods

Set up the bikeshop tables in your own database on the "ec2-35-176-213-89.eu-west-2.compute.amazonaws.com" database server using the the 2 SQL scripts provided in the chat. Use the v_customer_order_totals view to write a python function that takes the customer number as a parameter returns the order history for that customer

Write a python script to add 5% to the price of all products made by Trek. What is the sum total of all of the products made by Trek in the products catalogue?

The exercise should take approximately 45 minutes

Calling Stored Procedures using pyodbc

Calling stored procedures from pyodbc provides a more of a challenge. A well written stored procedure will return either a simple result set (in a very similar way that select statements do) or a simple value parameter.

A poorly designed stored procedure can return multiple different result sets from multiple select statements and also output parameters. Fortunately, pyodbc can cope with all of these eventualities.

The **cursor.nextset()** method moves the window onto the next result set coming back from the stored procedure and returns True, if there is indeed another result set. The new result set can then be accessed using one of the fetch methods we looked at earlier with simple selects. If there is no further result set calling `cursor.nextset()` returns False.



Be careful when calling this method as all rows remaining in the current result set are discarded and can only be accessed again by re executing the stored procedure

Handling input and output parameters with Stored Procedures using pyodbc

Input parameters with stored procedures are managed in the same way as we saw earlier with selects. We use one or more ? characters in the call and put substitution values in the parameters.

With output parameters we need to alter the SQL text of the call to declare a variable in TSQL, then reference it in the call to the procedure and then, after the call to the procedure completes, we select the value in a further statement and it is returned to the pyodbc cursor as another result set

```
5 myInputVariable = 'whatever I want - could be string/int/number whatever'
6
7 tSql = '''DECLARE @out1 int;
8         DECLARE @out2 varchar(50);
9         EXEC [dbo].[someDodgySP] ?, @out1 OUTPUT, @out2 OUTPUT;
10        SELECT @out1 AS my1stVarReturned;
11        SELECT @out2 AS my2ndVarReturned;'''
12
13 cursor.execute(tSql, myInputVariable)
14
15 # Use cursor.nextval to get any result sets from someDodgySP
16 # and then the my1stVarReturned and my2ndVarReturned variables
```

Exercise Ex 3

Selecting from a Stored Procedure with multiple result sets and Output Parameters

Use the bikeshop tables we created earlier

Write a customer summary stored procedure to summarise a customer's orders and their total spend with the company

The report should contain 3 sections.

- 1) A header with the customer's first name and last name, their email address and their city and state and phone number
- 2) The orders section should contain order_id the order date and the total order cost value
- 3) Finally return their total spend as an output parameter.

Format the results as a report in Python

Who are the top 3 biggest spenders with the bikeshop ?

Managing Type Conversions between Python and SQL

- Differences between the type system between Python and SQL databases
- Implicit type conversions performed by pyodbc
- Approaches to types that cannot be easily converted
 - Convert in-the-database to a type that can be handled by pyodbc
 - Perform output conversion in the pyodbc driver

Type Systems between Python and SQL Databases are Very Different

pyodbc does a great job of hiding the differences of approach between the two type systems



Python is Dynamically Typed

- You do not have to declare types on variables
- You can change the type of a variable after you use it
- Any type checking gets done at run-time



Database SQL implementations are Strongly Typed

- You must declare a type for a variable before you use it
- You cannot change a variables type without re-declaring it
- Type checking is done in a compile phase before run-time

Implicit Type conversions performed by pyodbc

- Since Python 3 string objects have been rendered in Unicode (utf-16 or utf-8)
- This has made the transferring of string data between Python and SQL databases fairly seamless (as we saw before)
- The recent SQL Server ODBC driver works out of the box with Python 3
- Some other databases and operating system platforms may need more configuration
- The connection methods `setEncoding()` and `setDecoding()` can be used to alter how the conversions are performed
- As always Michael Kleehammer's excellent documentation has many examples of common issues experienced on database platforms.
(<https://github.com/mkleehammer/pyodbc/wiki/Unicode>)

Implicit Type conversions performed by pyodbc passing data to the database

Python Datatype	Description	ODBC Datatype
None	null	varies (1)
bool	boolean	BIT
int	integer	SQL_BIGINT
float	floating point	SQL_DOUBLE
decimal.Decimal	decimal	SQL_NUMERIC
str	UTF-16LE (2)	SQL_VARCHAR or SQL_LONGVARCHAR (2)(3)
bytes, bytearray	binary	SQL_VARBINARY or SQL_LONGVARBINARY (3)
datetime.date	date	SQL_TYPE_DATE
datetime.time	time	SQL_TYPE_TIME
datetime.datetime	timestamp	SQL_TYPE_TIMESTAMP
uuid.UUID	UUID / GUID	SQL_GUID

1.If the driver supports it, [SQLDescribeParam](#) is used to determine the appropriate type. If it is not supported, SQL_VARCHAR is used.

2.The encoding and ODBC data type can be changed using [Connection.setencoding\(\)](#). See the [Unicode page](#) for more information.

3.[SQLGetTypeInfo](#) is used to determine when the LONG types are used. If it is not supported, 1MB is used.

Implicit Type conversions performed by pyodbc for data coming from the database

Description	ODBC Datatype	Python Datatype
NULL	any	None
bit	SQL_BIT	bool
integers	SQL_TINYINT, SQL_SMALLINT, SQL_INTEGER, SQL_BIGINT	int
floating point	SQL_REAL, SQL_FLOAT, SQL_DOUBLE	float
decimal, numeric	SQL_DECIMAL, SQL_NUMERIC	decimal.Decimal
1-byte text	SQL_CHAR	str via UTF-8 (1)
2-byte text	SQL_WCHAR	str via UTF-16LE (1)
binary	SQL_BINARY, SQL_VARBINARY	bytes
date	SQL_TYPE_DATE	datetime.date
time	SQL_TYPE_TIME	datetime.time
SQL Server time	SQL_SS_TIME2	datetime.time
timestamp	SQL_TIMESTAMP	datetime.datetime
UUID / GUID	SQL_GUID	str or uuid.UUID (2)
XML	SQL_XML	str via UTF-16LE (1)

- 1.The encoding can be changed using [Connection.setdecoding\(\)](#). See the [Unicode page](#) for more information.
- 2.The default is str. Setting pyodbc.native_uuid to True will cause them to be returned as uuid.UUID objects.

Approaches to use when there is no compatible type in pyodbc: Convert to a compatible type in the Database:

- In this demonstration we will show how we can use functions or methods in the database to render incompatible data types into string types that can be handled by the pyodbc driver.
- Here the EarthquakeData table has a spatial geography type that is specific to the SQL server database.
- We use the StAsText() method to convert it to a textual value

```
spatial_fail_eg.py X
spatial_fail_eg.py > ...

25
26 # We can use a python context manager with pyodbc.connect()
27 with pyodbc.connect(connectString) as cnxn:
28     # Create a cursor object from the connection
29     cursor = cnxn.cursor()
30     cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
31
32     # Select on a Geography type in Earthquakeinformation
33     Sql = '''SELECT EarthquakeID, Earthquakeinformation
34             FROM [djbpyodbc].[dbo].[EarthquakeData];'''
35
36     cursor.execute(Sql)
37
38     input("\nPress return to continue")
39     # Print out the cursor metadata for the last call
40     try:
41         allRows = cursor.fetchall()
42     except pyodbc.Error as pyodbcError:
43         print(f"Error: {str(pyodbcError)}")
44
45     input("\nPress return to continue")
46
47     # Define SQL BUT render point as a string
48     Sql = '''SELECT EarthquakeID, Earthquakeinformation.StAsText()
49             FROM [djbpyodbc].[dbo].[EarthquakeData];'''

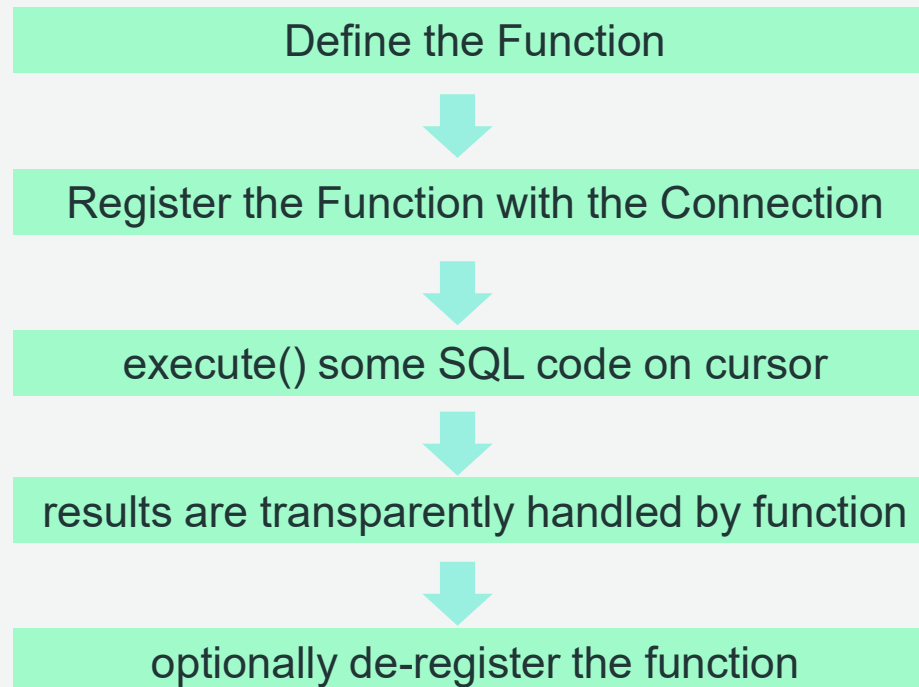
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

Error: ('ODBC SQL type -151 is not yet supported. column-index=1 type=-151', 'HY106')

Press return to continue

**Result description from cursor**
Name: EarthquakeID
Python type: <class 'int'>
```

Using Output Converter Functions



Approaches to use when there is no compatible type in pyodbc: Output converter functions in pyodbc

- In this demonstration we will show how we can use an output converter function that is run by the pyodbc driver to render incompatible data types into string types that we can use in our python code.
- Here the `dto_test` table has a datetime offset column, a type that is not handled by ODBC and the pyodbc module natively. In the demo we register a python function `handle_datetimeoffset()` to convert the datetime offset into a datetime data type we can work with in python.

```

spatial_fail_eg.py  dto_fail_eg.py X
dto_fail_eg.py > ...
15
16 def handle_datetimeoffset(dto_value):
17     # ref: https://github.com/mkleehammer/pyodbc/issues/134#issuecomment-281739794
18     tup = struct.unpack("<6hI2h", dto_value) # e.g., (2017, 3, 16, 10, 35, 18, 500000000, -6, 0)
19     return datetime(tup[0], tup[1], tup[2], tup[3], tup[4], tup[5], tup[6] // 1000,
20                     timezone(timedelta(hours=tup[7], minutes=tup[8])))
21
22 # config file with access creds to the database
23 with open('config.json') as f:
24     d = json.load(f)
25
26 hostname = d['hostname']
27 user = d['username']
28 psw = d['psw']
29 database = 'djbpodbc'
30
31 connectionString = 'DRIVER={ODBC Driver 18 for SQL Server};SERVER='+\
32 hostname+';DATABASE='+database+';UID='+\
33 user+';PWD='+psw+';TrustServerCertificate=yes'
34
35 # We can use a python context manager with pyodbc.connect()
36 with pyodbc.connect(connectionString) as cnxn:
37     # Create a cursor object from the connection
38     cursor = cnxn.cursor()
39     cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
40
41 # Select on a Date time Offset type in dto_test
42 Sql = '''SELECT id, dto_col
43         FROM dbo.dto_test'''
44
45 PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL
Python type: <class 'str'>
Nullable: True

(1, datetime.datetime(2017, 3, 16, 10, 35, 18, 500000, tzinfo=datetime.timezone(datetime.timedelta(days=-1, seconds=64800))))

PS Microsoft.PowerShell.Core\FileSystem::\\ws1$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python>

```


In-Database Conversion or Output Convertor Functions

How to choose which ?

In Database Type Conversion

- In-Database conversion can be made to work on both values passing to and from the database
- In-Database conversion is generally easier to code
- In database conversion requires that you are able to make changes to the database and that is not always possible

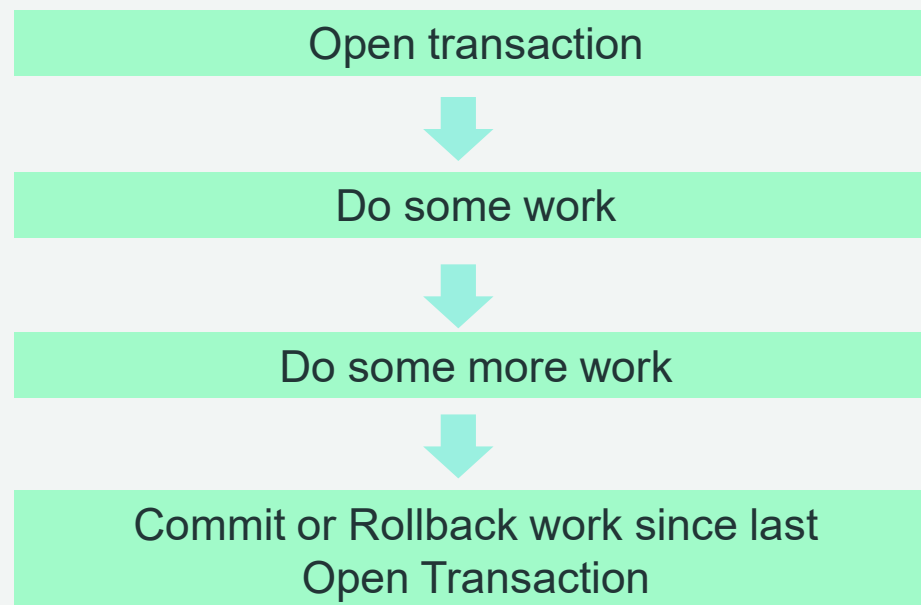
Output Convertor Functions

- Only convert data coming from the database to pyodbc, passing to the database
- Performed in pyodbc – in Python
- Might be your only option if you cannot make changes to the database

SQL Databases bundle operations together in Transactions

- Most SQL Database systems support Database Transactions
- Transactions are a way in which databases can bundle together related operations
- Consider a bank needing to transfer money from one account to another
They might perform this as 2 updates to the accounts table
1st update reduces balance on account 1 by the amount of money to be transferred
2nd update increases balance on account 2 by the amount of money to be transferred
- The bank would want either both updates to occur OR neither
If either of the updates occur on their own – then the bank has a problem
- Transactions allow the grouping of these changes together so they all succeed OR they all fail
- You may hear of people referring to **ACID** transaction properties
 - Atomic** – the transaction is an indivisible unit of work that all happens or none of it does
 - Consistent** – the database system moves from one state where whole transactions have been applied to another – with no visible state where parts of transactions have been applied
 - Isolated** – while one transaction of work is being applied it doesn't affect any others in progress
 - Durable** – when a transaction has succeeded it is recorded securely in the database

Setup and commit for transactions



How pyodbc Manages Transactions

- The main mechanism that pyodbc uses to control how it handles transactions is the **connection** autocommit property
- You can set autocommit to True – in which case pyodbc will commit each statement as it is submitted (and you effectively have no control over transactions in pyodbc)
- If you set autocommit to False (which is the default) then
 - when you open a cursor there is an implicit “begin transaction”
 - you can perform some work in the cursor in a bundle
 - you can issue a commit statement to save all of your work bundle to the database
 - if you detect an error you can issue a rollback statement to rollback all of the bundle of work



There is a feature of the way pyodbc handles transactions that could catch you out. Transactions are handled at the **connection** level but there are `cursor.commit()` and `cursor.rollback()` methods. These cursor method affect database transactions at the connection level and they are functionally identical to the connection commit and rollback methods. If you have several cursors generated from one connection and you issue a commit – the work in ALL of the cursors gets committed (or rolled back if you issue a rollback). This might not be what you were expecting!

Demonstration with connection autocommit True and False

- Here we show the effect on a session when the connection autocommit is set to True and when autocommit is set to False

```

25 # We can use a python context manager with pyodbc.connect()
26 with pyodbc.connect(connectString) as cnxn:
27     # Create a cursor object from the connection
28     cursor = cnxn.cursor()
29     cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
30     #autocommit False by default
31     print(f"Autocommit is {str(cnxn.autocommit)}\n")
32
33     # Create a new table with special offer items in it
34     cr_tab = '''CREATE TABLE [production].[promo_products](
35     [promo_id] [int] IDENTITY(1,1) NOT NULL,
36     [product_id] [int] NOT NULL,
37     [product_name] [varchar](255) NOT NULL,
38     [brand_id] [int] NOT NULL,
39     [category_id] [int] NOT NULL,
40     [model_year] [smallint] NOT NULL,
41     [list_price] [decimal](10, 2) NOT NULL,
42     PRIMARY KEY CLUSTERED
43     (

```

PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

About to creat table Press return to continue

About to poulate table Press return to continue

About to Rollback Press return to continue

About to DROP table Press return to continue

About to Rollback Press return to continue

Press return to End

PS Microsoft.PowerShell.Core\FileSystem::\\ws1\$\\Ubuntu-20.04\\home\\djbadm\\src\\SQL_DB_Access_Python>

Setting autocommit to True – one last Gotcha !



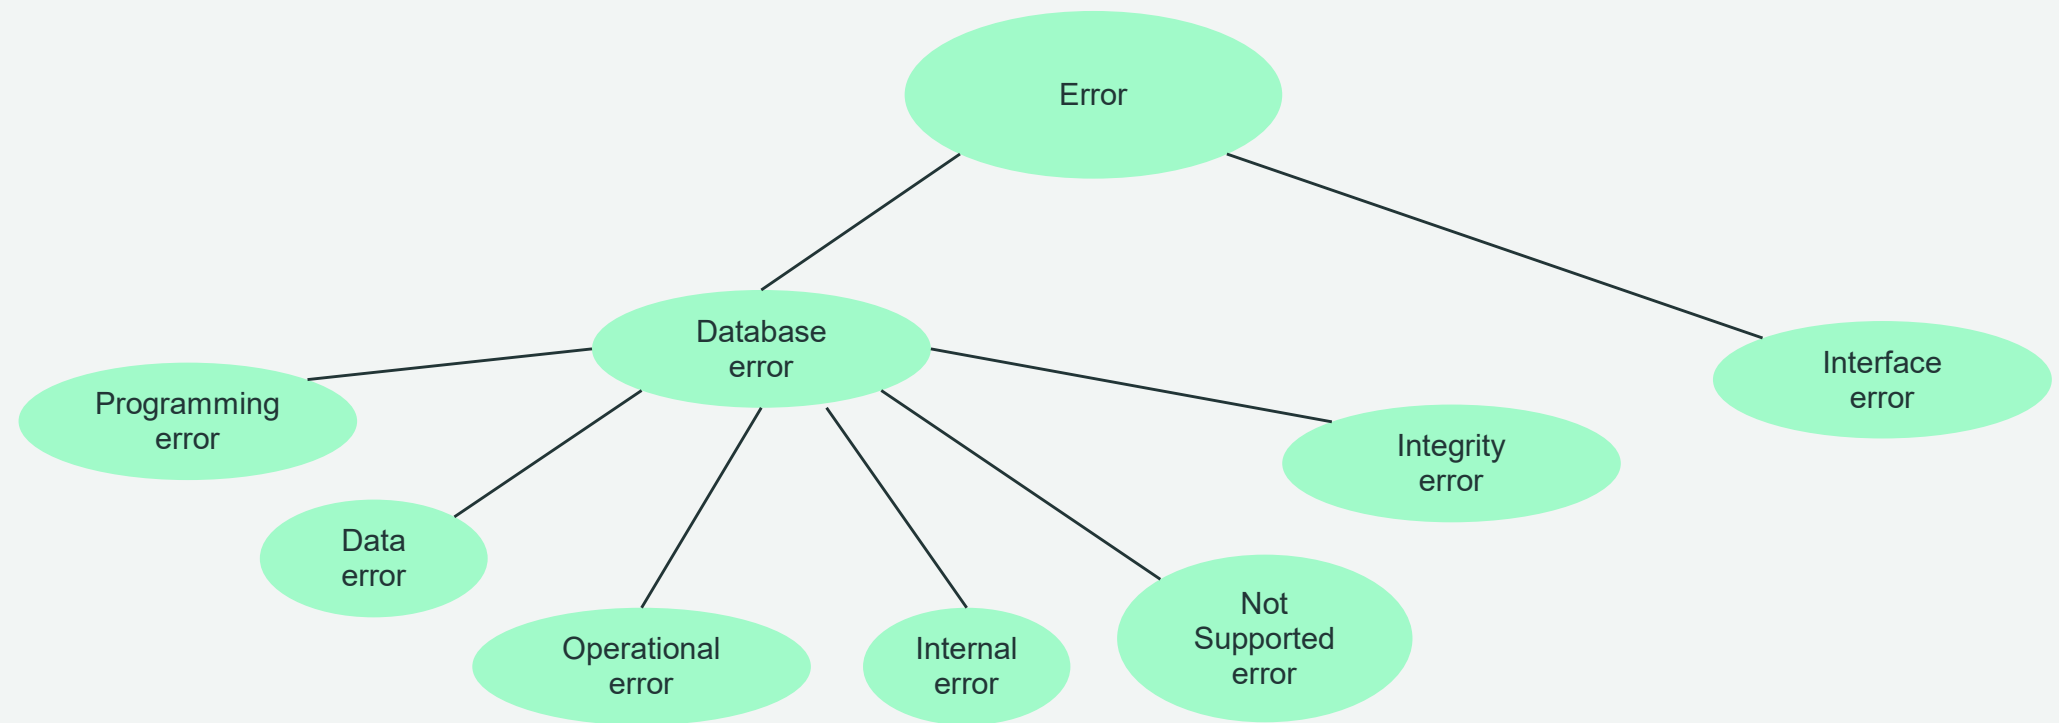
If you set the autocommit connection property to True there is one further thing to be aware of. If you use the `executemany()` method on a cursor and the connection `fast_executemany` is set to False (as is the default). Then each parameter set is sent as a separate SQL statement (with its own implicit commit as autocommit is set to True)

If a database failure occurs during the operation – it can be very difficult to tell what has been committed and what part of the batch has not.

Specifically, for Microsoft SQL Server, if you have to use the connection autocommit to true then you should also set `fast_executemany` to True in the connection

Handling Errors in pyodbc

The pyodbc driver uses the standard pattern of raising exceptions for various errors that may occur when connecting to and interacting with the database. The hierarchy of error classes is as follows



Handling Errors try except finally pattern

Errors can be captured and handled using the standard Python try, except catch blocks with optional finally blocks to clear up after failures

The pattern is generally as follow
try:

 some database operation

except pyodbc.ErrorType1 as pyodbcerror:
 take some suitable action

except pyodbc.ErrorType2 as pyodbcerror:
 take some suitable action

except pyodbc.ErrorType2 as pyodbcerror:
 take some suitable action

finally:
 optional clean up (if you have to terminate etc)

Might want to have multiple except blocks to cope with multiple classes of error

Handling Errors Demo

We will look at how we use the try, catch blocks to handle errors coming back from the database

```
demo_local_error_handling.py > ...
19 database = 'dikesnop'
20
21 connectionString = 'DRIVER={ODBC Driver 18 for SQL Server};SERVER='+\
22 hostname+';DATABASE='+database+';UID='+\
23 user+';PWD='+psw+';TrustServerCertificate=yes'
24
25 input(" shutdown the local SQL database!")
26 try:
27     cnxn = pyodbc.connect(connectionString)
28 except pyodbc.OperationalError as pyodbcerror:
29     print(f"Error {str(pyodbcerror)}")
30     input("\nStart the database now Press return to continue")
31
32 # startup the local database and get the connection
33
34 cnxn = pyodbc.connect(connectionString)
35 # Create a cursor object from the connection
36 cursor = cnxn.cursor()
37 cursor.setinputsizes([(pyodbc.SQL_WVARCHAR, 50, 0)])
38 #autocommit False by default
39 print(f"Autocommit is {str(cnxn.autocommit)}\n")
40 input("\nEstablished a connection Press return to run disconnect demo\n")
41 print("Database Disconnect Demo")
42 sql = "WAITFOR DELAY '00:01:30';"
43 try:
44     cursor.execute(sql)
45     rows = cursor.fetchall()
```

Exercise Ex4

Handling Errors and managing transactions

Work in your pods and use the bikestore tables we created earlier

Write a customer transaction simulation: The simulation should behave like an order processing system.

Create new orders for existing goods in the production.products tables. The orders can come from existing customers (to make things a little simpler) and are for existing products.

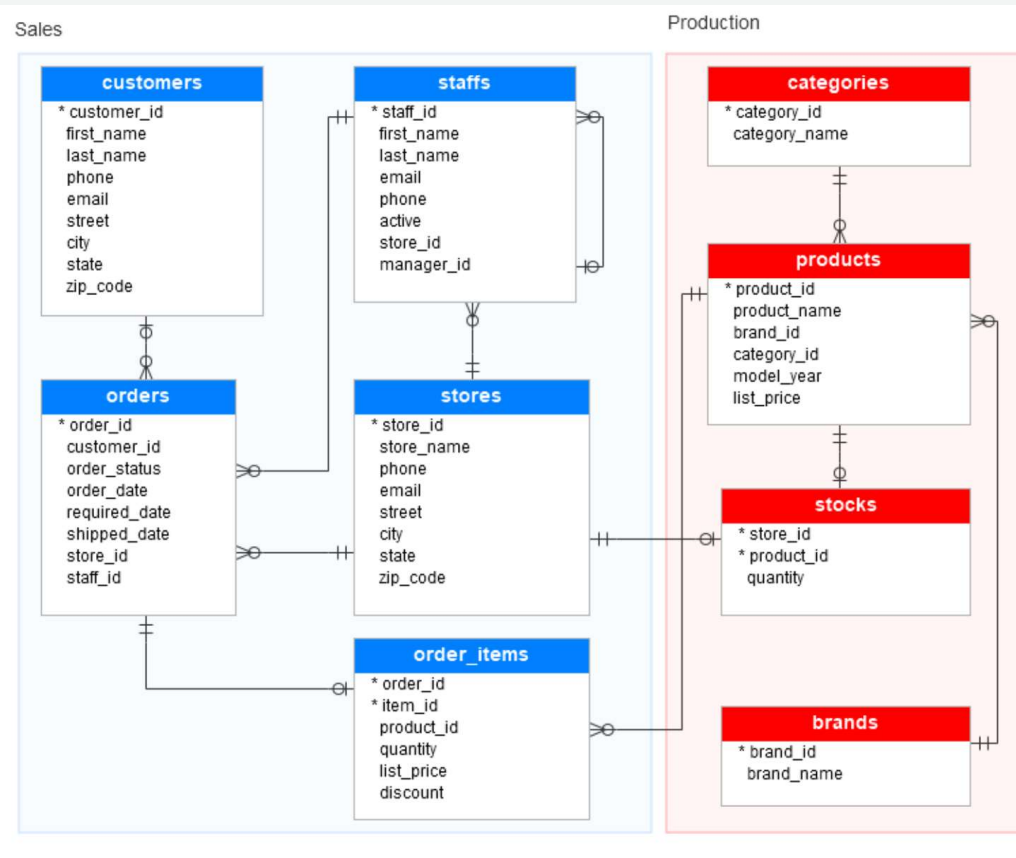
When new orders come in you should update the number of items from the stocks table and you should allocate a salesperson and a store to the order.

You should record the new orders in the sales.orders and order items tables.

Use the order_status field to signify whether an order has been placed , fulfilled , cancelled etc – you may have to create a reference table for this.

Simulate some orders failing due to lack of stock, payment declined, etc . Simulate failures and use transaction rollbacks to tidy up after failing transactions.

Decide as a team which parts of the solution to perform in python and which in SQL



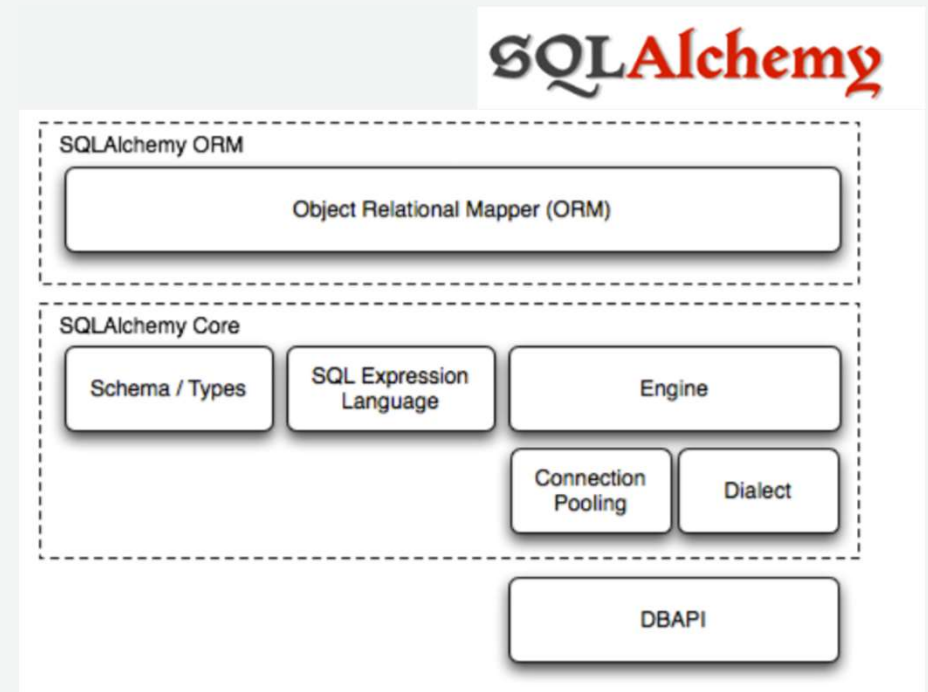
SQLAlchemy

SQLAlchemy is a higher abstracted API than pyodbc and it sits on top of the DBI 2.0 specification.

It and provides more insulation against differing Database vendors than is possible with pyodbc. It has dialects to cope with the differing features of SQL databases. It has a metadata language to describe database objects and queries to facilitate this.

SQLAlchemy provides connection pooling and allows for larger applications requiring multiple concurrent connections to manage their database access more easily

SQLAlchemy engine can be used by Pandas to rapidly load dataframes from database tables and vice-versa.



Further Reading

Michael Kleehammer's wiki pages <https://github.com/mkleehammer/pyodbc/wiki>

SQLAlchemy – Python relational toolkit and Object Relational Mapper
<https://www.sqlalchemy.org>

kubrick

Questions

11 April 2023

